
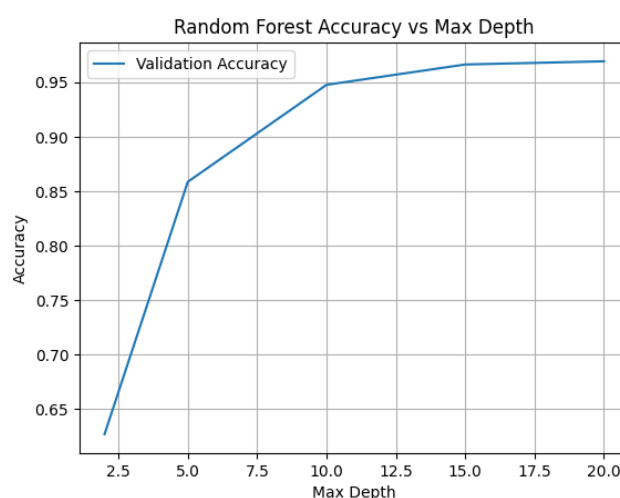



Evaluating model's performance for SVM, RF and LSTM using MNIST dataset

Random Forest

```
{  
  "n_estimators": range(50, 1050, 50),  
  "max_depth": {2, 5, 10, 15, 20, None},  
  "min_samples_split": range(2, 42, 2)  
}
```



Seeing that the best number of **estimators** is 650, best number of **max_depth** is 20 and best number of **min_samples_split** is 2. In combination, the resulting test accuracies yields


Validation accuracy: **0.971**
Validation Precision: **0.9710116403292254**
Validation Recall: **0.971**
Validation F1 Score: **0.9709933654164185**

```
Cross Validation Scores: [0.96866667 0.96708333 0.96533333 0.96525
0.97183333]
Mean Cross Validation Score: 0.9676333333333333
```

As the number of estimators of 650 is quite high, and may result in higher computational cost with diminishing returns, I choose 150 to be the best combination of accuracy and efficiency. Indeed, the test yields

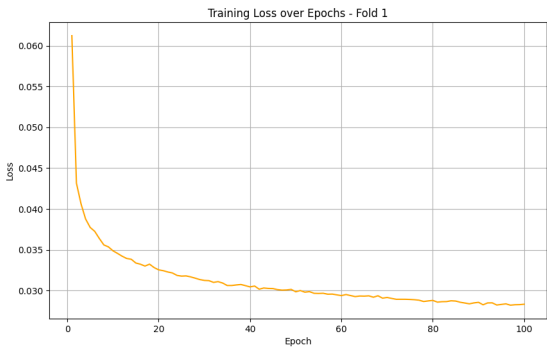

```
Validation accuracy: 0.9696
Validation Precision: 0.9695990508409351
Validation Recall: 0.9696
Validation F1 Score: 0.9695861419399533
Cross Validation Scores: [0.96675    0.96475    0.9635    0.96316667
0.971    ]
Mean Cross Validation Score: 0.9658333333333333
```







which only differs a mere .002 percentage point (this could be amount to the noise in the progress of training, but after 10 experiments, the results seems consistent)



SVM

```
{  
  "n_epochs": 100  
}
```








Fold 1 Results:
Accuracy: 0.9107
Precision: 0.9113
Recall: 0.9107
F1 Score: 0.9108

Fold 2 Results:
Accuracy: 0.9077
Precision: 0.9079
Recall: 0.9077
F1 Score: 0.9074

Fold 3 Results:
Accuracy: 0.9094
Precision: 0.9096
Recall: 0.9094
F1 Score: 0.9092

Average Results across all folds:

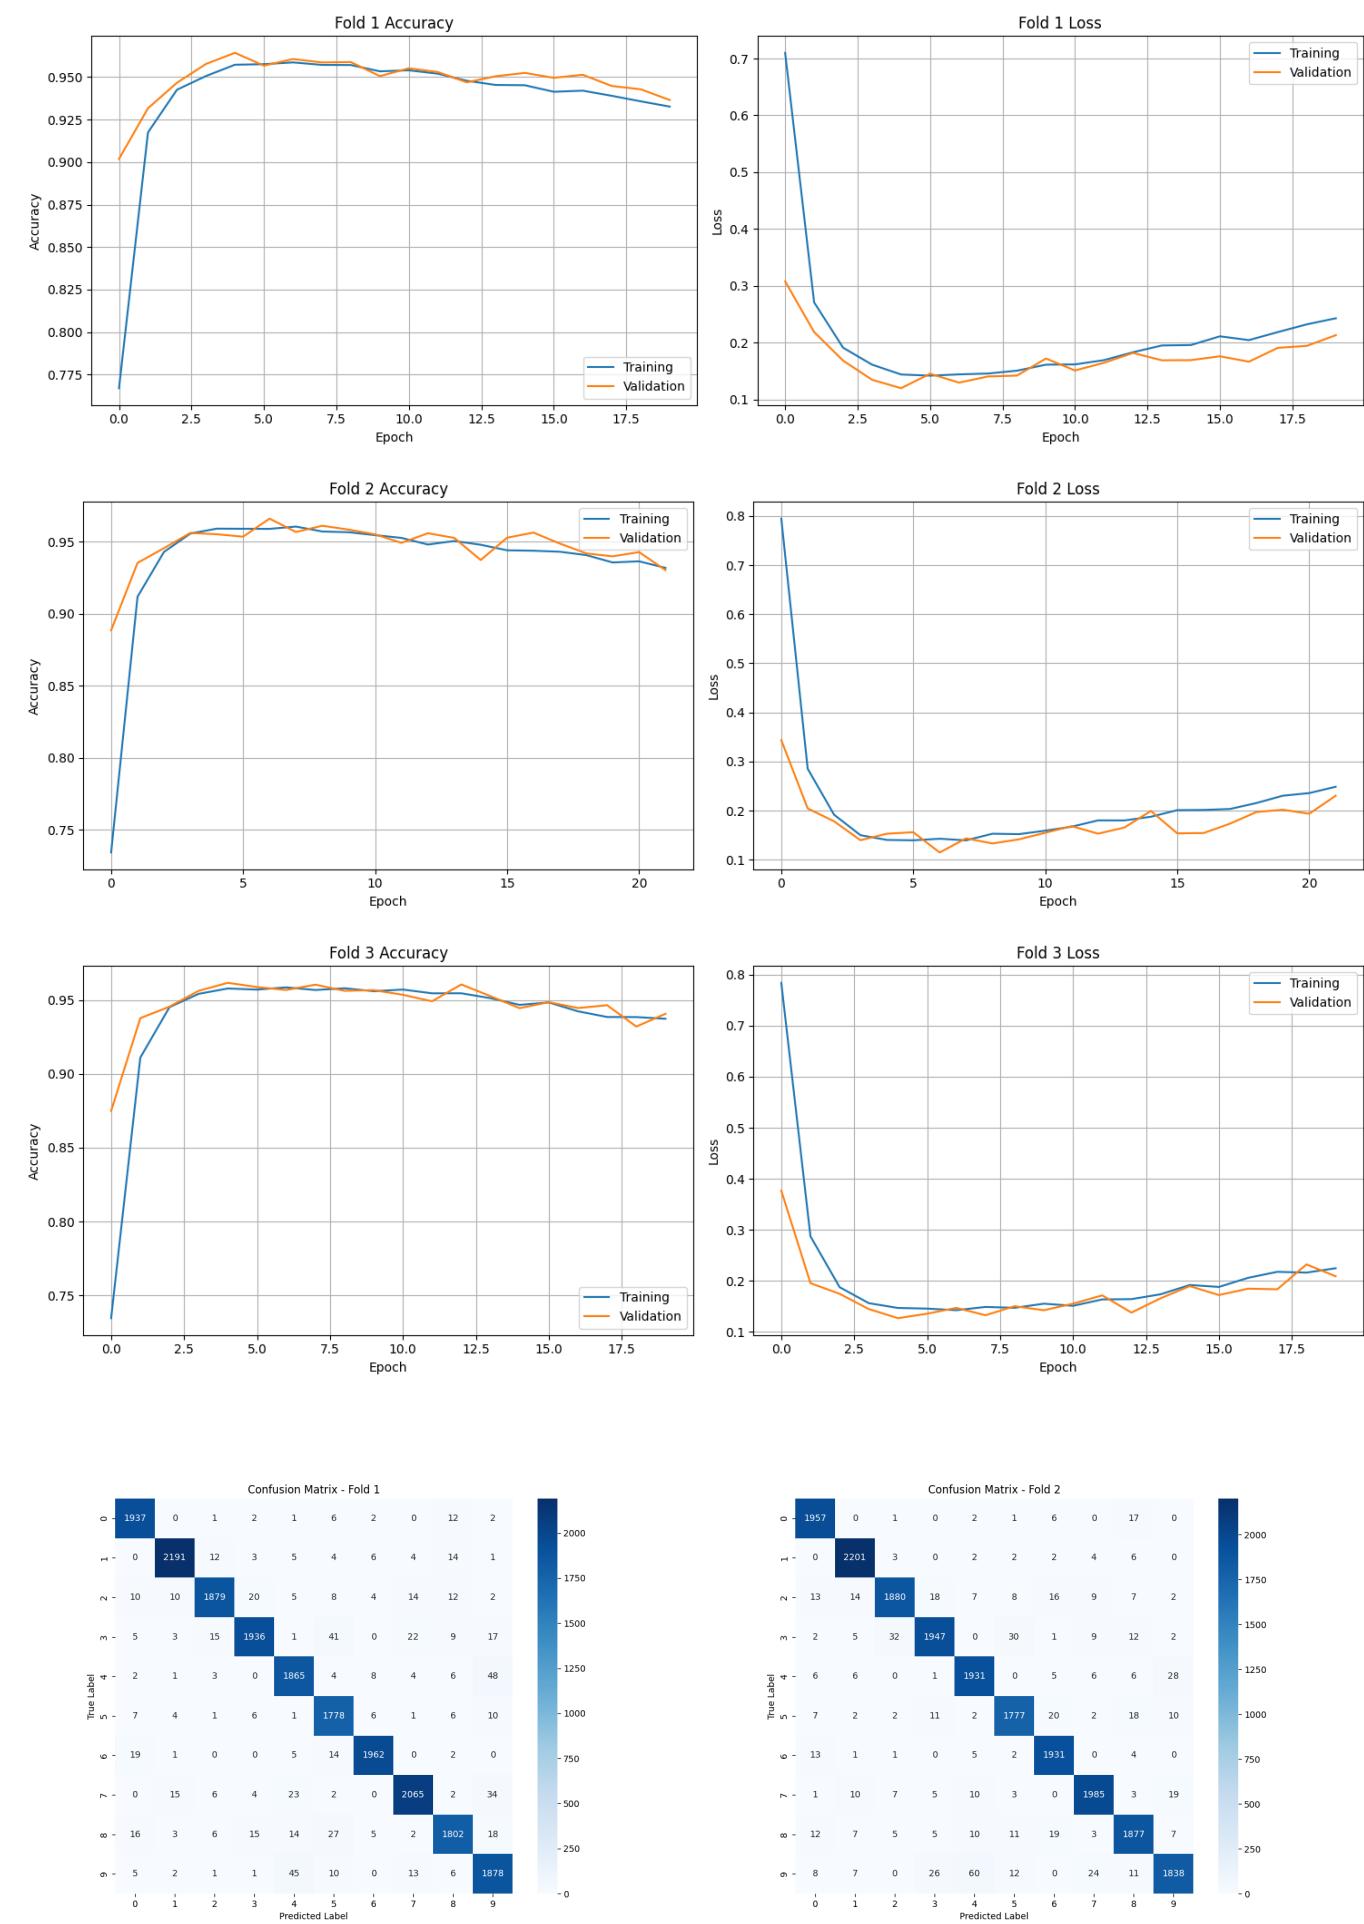
Accuracy: 0.9093
Precision: 0.9096
Recall: 0.9093
F1 Score: 0.9092




LSTM

```
{  
  "n_epochs": 100  
}
```

Below is the result of training v validation accuracy through each of the 100 epochs of LSTM





Early stopping triggered after epoch 20

Fold 1 Final Results:
Accuracy: 0.9647
Precision: 0.9649
Recall: 0.9647
F1 Score: 0.9647

Early stopping triggered after epoch 22

Fold 2 Final Results:
Accuracy: 0.9662
Precision: 0.9662
Recall: 0.9662
F1 Score: 0.9661

Early stopping triggered after epoch 20

Fold 3 Final Results:
Accuracy: 0.9627
Precision: 0.9630
Recall: 0.9627
F1 Score: 0.9627

Average Results across all folds:
Accuracy: 0.9645
Precision: 0.9647
Recall: 0.9645
F1: 0.9645